

Response to Reviewers
BMC Genomics submission e0f004a9-2cfd-4839-8061-e650d2bb8143
“A Multi-Tissue Transcriptomic Atlas of Porcine Development Identifies Conserved
Molecular Programs with Humans”

Tianyuan Liu Rujing Lei Ilyas M. Khan Peter S. Theobald*

May 11, 2026

Dear Dr Kanade and Dr Harrison,

We thank you and the three reviewers for the constructive and thorough feedback on our manuscript. We have addressed every point raised, performed the requested additional analyses, and revised the manuscript accordingly. The most substantive changes are:

1. A dedicated batch-effect evaluation across all five tissues (new Supplementary Figure S4 and Supplementary Table S6), colouring PCA and UMAP embeddings by **BioProject** as recommended by Reviewer 3 (R3.1).
2. Tissue-specific top-feature tables for brain, liver, lung, and blood (R3.4) (new Supplementary Figure S5 and Supplementary Table S7).
3. An independent biological validation of stage boundaries using a literature-curated marker-gene panel (R3.3) (new Supplementary Figure S6).
4. Cross-species concordance extended from skeletal muscle alone to brain and liver, using the Cardoso-Moreira *et al.* (2019) developmental atlas as recommended by Reviewer 1 (R1.2) (new Supplementary Figure S7 and Supplementary Table S8). Note: human lung is not present in the Cardoso-Moreira atlas; that comparison cannot be performed with this dataset and is flagged as a future-work item.
5. A within-stage heterogeneity/sensitivity analysis stratifying by breed, sex, and BioProject (R2.1) (new Supplementary Figure S8 and Supplementary Table S9).
6. Explicit formulae for the Frank & Hall reduction and an accurate description of how monotonicity is enforced *post hoc* (not via the LightGBM `monotone_constraints` parameter), per R2.2.
7. One-to-one Ensembl Compara ortholog filter and full human dataset preprocessing pipeline now documented explicitly in the Methods (R3.2).
8. Strengthened limitations: postnatal-only window (R1.1), small human cohort and infant-vs-adult gap (R3.m2), and chronological-age-bin heterogeneity (R2.1).
9. Editorial audit of overstated claims: every quantitative statement is now scoped to the dataset it derives from.

A point-by-point response follows. Reviewer comments are shown verbatim in shaded boxes; our responses, with cross-references to the revised manuscript, appear beneath each box.

*Corresponding author: TheobaldPS@cardiff.ac.uk

Response to the Editor

Editor (Paul Harrison)

There are several methodological issues that the reviewers point out.

Response: We thank the editor for the careful handling of our manuscript. Every methodological point raised by the three reviewers has been addressed below. In line with the editor’s instruction to “ensure the results are accurately reported, any overstated conclusions are rewritten and the limitations of the work fully explained”, we conducted a full audit of the manuscript text and revised every overstated claim. Specifically: (i) all cross-species claims are now qualified by the small human cohort size ($n = 11$); (ii) batch-effect robustness statements are bounded by the BioProject coverage we could test (detailed in new Supp. Fig. S4); (iii) the “entire postnatal lifespan” phrasing in Result 2.1 has been replaced with “the postnatal lifespan from 0–730 days”, and a new paragraph (R1.1) explicitly excludes prenatal development from our scope.

Manuscript changes: Throughout `paper.tex`; explicit additions detailed in each R-point below.

Reviewer 1

R1.1 — Postnatal-only scope

For Result 2.1 *A comprehensive transcriptomic atlas of porcine development*. The authors also described “The dataset comprises 1,924 samples spanning five major tissues (brain, liver, muscle, lung, and blood) and covering the entire postnatal lifespan”. But, as we all known, the development of any mammals should include prenatal and postnatal stages. However, the authors only collected bulk RNA-seq data from porcine postnatal stages. So, the authors should explain their limitations.

Response: We agree completely. The original phrasing “the entire postnatal lifespan” implicitly suggested whole-of-development coverage, which we did not intend. We have (i) replaced the offending sentence in Results 2.1 to make the postnatal-only window explicit (“covering the postnatal lifespan from the day of birth to beyond 365 days of age; prenatal (embryonic and fetal) development is outside the scope of the present atlas, since the PigGTEx resource does not include systematic, multi-tissue prenatal sampling and addressing this developmental window would require dedicated prenatal sample collection”), and (ii) added an explicit Limitations item (now the *first* limitation enumerated) clarifying that no claim is made about prenatal trajectories and that integration with single-cell atlases of porcine organogenesis is the natural complement to our postnatal staging. We thank the reviewer for catching this important framing issue.

Manuscript changes: `paper/paper.tex` Result 2.1 paragraph (the sentence beginning “The dataset comprises 1,924 samples...”) and the first paragraph of the Limitations subsection in the Discussion; `thesis/MPhilThesis-Latex-Template/5_Discussion/5_Discussion.tex` new “Postnatal-Only Atlas Scope” subsubsection.

R1.2 — Cross-species expansion (Cardoso-Moreira 2019)

For Result 2.4 *Evolutionary conservation with human development*. The authors only performed a cross-species analysis using the skeletal muscles between human and pig. However, a previous study (DOI: 10.1038/s41586-019-1338-5) has generated bulk RNA-seq data of brain, liver, and lung in human, macaque, mouse, rat, rabbit, opossum and chicken across their developmental time points from early organogenesis to adulthood.

Response: We are grateful for this excellent suggestion—the Cardoso-Moreira *et al.* 2019 atlas[1] is an appropriate data source to extend cross-species concordance beyond muscle. We have done so for the available tissues and report the results honestly, including one substantive caveat the reviewer should be aware of.

Data caveat first. The Cardoso-Moreira 2019 human RPKM matrix (`Human.RPKM.tsv`, E-MTAB-6814) contains seven organs: Brain, Cerebellum, Heart, Kidney, Liver, Ovary, Testis. *It does not include lung.* Of the three reviewer-requested tissues, Brain and Liver are present and have been analysed; Lung is not available and we cannot infer human postnatal lung conservation from this atlas. We state this explicitly in the new Results 2.4 paragraph and in Supplementary Fig. S7 / Supplementary Table S8.

Method (also addresses R3.2). We applied exactly the same pipeline used for muscle: strict one-to-one Ensembl Compara orthologs ($n = 7,485$ validated pairs at the Ensembl Compara REST homology endpoint; see also the R3.2 response); $\log_2(\text{Infant/Adult})$ in both species using $\log_2(x + 0.01)$ on raw RPKM/TPM (matching the original pig pipeline); $\text{FDR} < 0.10$ and $|\log_2 \text{FC}| > 0.5$ in both species (“strict” filter). Because the per-tissue postnatal sample size in Cardoso-Moreira is small (Brain: 7 young + 9 adult human samples; Liver: 4 young + 6 adult), per-gene human FDR is underpowered, so we additionally report a pig-anchored variant using only the pig FDR filter; this is openly flagged as exploratory.

Findings (Supplementary Fig. S7; Supplementary Table S8).

- Brain (strict, $n = 272$): Pearson $r = 0.29$, 95% CI $[0.17, 0.41]$, $p = 1.6 \times 10^{-6}$, directional concordance 85.3%.
- Brain (pig-anchored, $n = 2,193$): $r = 0.21$, $p = 2.7 \times 10^{-22}$.
- Liver (strict, $n = 0$): no orthologs pass the strict bidirectional filter (consistent with the small human cohort).
- Liver (pig-anchored, $n = 857$): $r = 0.21$, $p = 3.0 \times 10^{-10}$, directional concordance 47.3%—*below the 50% chance baseline.*

The brain strict result is statistically significant but its top genes by pig $\log_2 \text{FC}$ are dominated by canonical skeletal-muscle markers (*TNNC2*, *SMPX*, *MYBPHL*, *CHRNA1*, *LRRC2*), whose median expression in pig “brain” is near zero with high-TPM outliers in a small number of samples. We therefore interpret the brain r as reflecting cross-species concordance of muscle developmental dynamics that leak into the brain expression matrix—consistent with variable muscle-tissue contamination of a subset of pig brain biopsies in PigGTEx—rather than as a clean estimate of brain-specific developmental conservation. The liver pig-anchored estimate falls below the 50% chance baseline for direction, so we report it as a null result rather than as a smaller but compatible signal. We have updated the Results 2.4 paragraph and Supplementary Table S8 caption to state both points explicitly.

Honest framing. Lung is genuinely unavailable in the Cardoso-Moreira atlas; brain is a partial positive result with a biological-leakage caveat; liver is a null result given the available human cohort. The cleanest cross-species finding remains the original muscle analysis with a matched human muscle cohort. Obtaining clean per-tissue conservation estimates for brain, liver, and lung will require dedicated human postnatal cohorts with matched dissection rigour, and we say so plainly in the Discussion.

The full analysis is reproducible from `review/analyses/cross_species_all_tissues.py`, with results in `review/analyses/results/cross_species_all_tissues.json` and gene-level data in `review/analyses/results/cross_species_per_tissue_genes.csv`.

Manuscript changes: New “Cross-species expansion to additional tissues” paragraph at the end of Results 2.4 in `paper/paper.tex`; new Supplementary Fig. S7 and Supplementary Table S8; `thesis/.../5_Discussion/5_Discussion` already references Supplementary Fig. S7 and Supplementary Table S8; ortholog filter and human preprocessing pipeline added to the Methods section (linked to R3.2 response).

Reviewer 2

R2.1 — Within-stage heterogeneity (breed/sex/project)

The study defines five developmental stages based on chronological age bins (e.g., infant: 0–20 days). While these bins align with known physiological milestones, substantial heterogeneity likely exists within each stage due to genetic, environmental, or management factors. The manuscript should discuss how this heterogeneity is addressed by the model or acknowledge it as a limitation. For instance, could the model’s performance be influenced by uneven sampling across sub-categories (e.g., different breeds, sexes, or rearing conditions)? A sensitivity analysis or discussion of this variability would strengthen the robustness claims.

Response: The reviewer is correct on both counts: substantive within-stage heterogeneity exists, and it can in principle bias the marginal balanced accuracy (BA) figures we report. We have now added a dedicated **stratification audit and subgroup-sensitivity analysis** (new Supplementary Fig. S8 and Supplementary Table S9; analysis script `review/analyses/heterogeneity_sensitivity.py`). The analysis replays the same 5-fold stratified split (`StratifiedKFold`, `seed = 42`) and refits each fold with the per-fold hyperparameters already stored in `model_outputs/{Tissue}_results.json`, so the per-sample CV predictions are bit-identical to the published metrics (fold BAs match the stored values to four decimal places).

Heterogeneity is real and uneven across stages. In muscle the dominant breed identity changes between bins (Infant: 46% Landrace; Early childhood: 62% Duroc; Post-pubertal: 42% Duroc $\times$ Pietrain); the Early childhood subset is 72% female (55/76); Pre-pubertal is 54% male (49/90); the Adult stage is represented by only five samples of unknown breed. In liver, per-stage breed dominance spans 29–62%. Brain Infants are 96% a single PIC cross. Blood Post-pubertal is 100% “Unknown” breed. These imbalances are tabulated in `heterogeneity_stratification.csv`.

Subgroup performance varies—sometimes substantially. For muscle, per-breed BA ($n \geq 10$) spans 0.61 (Large white, $n = 68$) to 1.00 (Yorkshire, $n = 12$); per-sex BA spans 0.70 (Unknown-sex pooled, $n = 118$) to 0.95 (Female, $n = 274$); per-BioProject BA spans 0.50 (PRJEB28465, $n = 48$, only two stages represented) to 1.00 (PRJNA341938, $n = 21$). A within-stage label permutation test (1,000 shuffles) confirms the spread is larger than expected by chance for muscle/Sex ($p = 0.005$), muscle/BioProject ($p = 0.015$), and liver/Breed ($p = 0.008$); muscle/Breed sits at the boundary of nominal significance ($p = 0.054$) and is therefore reported as borderline rather than excluded; Brain/Sex, Liver/Sex, and Lung/Sex are consistent with chance ($p \geq 0.23$).

Honest implication. The marginal point estimates remain within the range we report (Muscle 0.89, Liver 0.92, Brain 0.81, Lung 0.89, Blood 0.96), and no subgroup with $n \geq 50$ falls below 0.70 for muscle or liver. The classifier should therefore be interpreted as calibrated to the breed/sex composition of PigGTEx; deployment to cohorts dominated by under-represented breeds (notably Large white for muscle, where BA = 0.61) warrants re-calibration on the new cohort. We have added this point both to the Discussion and to the web-app input documentation.

Manuscript changes: New Supplementary Fig. S8 and Supplementary Table S9; `paper/paper.tex` expanded Limitations item “Sixth...” to quantify within-stage heterogeneity; `thesis/.../5_Discussion/5_Discussion.tex` new “Within-Stage Demographic Heterogeneity” subsubsection.

R2.2 — Frank & Hall formula & monotone constraints

The Methods section describes the use of the Frank & Hall reduction. For complete transparency, it would be helpful to specify how the final stage probability is computed from the K-1 binary sub-model predictions (e.g., using the formula $\Pr(y=c) = \Pr(y > c-1)$)

- $\Pr(y > c)$). Additionally, a brief note on how the monotone constraints were applied in LightGBM to enforce the ordinal relationship of the cumulative probabilities would be useful for practitioners seeking to implement similar models.

Response: We thank the reviewer for these methodological clarification requests; the points are well-taken and have improved the transparency of the manuscript. We have expanded the Methods accordingly in two ways. First, the recovery of class probabilities from cumulative predictions is now stated explicitly with the exact formula the reviewer suggested:

$$\Pr(y = c \mid x) = \Pr(y > c - 1 \mid x) - \Pr(y > c \mid x),$$

under the boundary conventions $\Pr(y > -1 \mid x) = 1$ and $\Pr(y > K-1 \mid x) = 0$, with explicit edge cases for $c = 0$ and $c = K - 1$. Second, we have corrected a misleading statement about monotone constraints. Our implementation does *not* use LightGBM’s `monotone_constraints` parameter (which constrains the sign of the partial derivative of a tree ensemble with respect to a specific input feature, and is therefore unrelated to the ordinal output structure). Instead, monotonicity of the cumulative probabilities is enforced through two deterministic post-hoc projections at inference time: (i) sequential clipping of the cumulative probabilities so that $\widehat{\Pr}(y > k) \leftarrow \min(\widehat{\Pr}(y > k), \widehat{\Pr}(y > k - 1))$ for $k \geq 1$, restoring the non-increasing ordering required by the differencing identity above, and (ii) projection of the resulting class probabilities onto the unimodal simplex by propagating non-increasing values away from the predicted mode. Both projections are deterministic, applied uniformly during training and inference, and are implemented in the `OrdinalLightGBM.predict_proba` method of the released codebase. We hope this is a useful template for practitioners.

*Manuscript changes: **paper/paper.tex** Methods section “Machine learning framework” (Equation (1) is unchanged; the surrounding paragraphs have been expanded with explicit boundary conditions and an accurate description of the post-hoc monotonicity projections); **thesis/MPhilThesis-Latex-Template/3_Methodology/3_Methodology** “Monotonic Constraint Post-Processing” subsection updated to match.*

R2.3 — Web app input requirements / preprocessing guide

The interactive web application is a valuable resource. However, the manuscript should clarify the specific input requirements for users. Since the model was trained on 31,908 genes, what are the recommendations for users who might have RNA-seq data with different gene annotations or microarray data? Providing guidelines on data preprocessing or feature mapping in the manuscript or supplementary materials would increase the tool’s usability.

Response: Thank you for this practical suggestion—this is exactly the information users need to use the application safely. We have addressed the request at two levels.

In the manuscript. A new paragraph in the Discussion (“Input requirements and preprocessing recommendations for the web application”) now spells out:

- The expected input format (CSV or TSV; samples $\times$ genes or the transposed orientation, auto-detected) and the expected gene identifier system (*Sus scrofa* Ensembl ENSSCG; version suffixes are stripped).
- The expected expression units (raw TPM, with internal $\log_2(x + 1)$; a checkbox allows already-log-transformed input).
- The behaviour when genes are missing from the user’s matrix (zero-imputed); the gene-coverage fraction is reported in-app and we recommend $\geq 80\%$ before trusting the prediction.
- Concrete recommendations for the three off-distribution scenarios the reviewer mentions: (i) non-Ensembl annotations (lift to ENSSCG via `mygene.info` or BioMart); (ii) microarray

data (probe-set → gene ID aggregation, mean/median across probes, upload as log-transformed values with the caveat that dynamic range differs); (iii) cohorts dominated by under-represented breeds (notably Large white for muscle), where re-calibration on a small in-house cohort is recommended (this point is shared with R2.1).

In the application. A second expandable help panel (“Using data with different annotations or microarray data”) has been added to the application’s upload page, presenting the same guidance points adjacent to the file uploader so users encounter them *before* they upload data. The panel also now covers single-cell RNA-seq (pseudo-bulking guidance) and out-of-distribution cohorts (re-calibration guidance).

Manuscript changes: `paper/paper.tex` new “Input requirements and preprocessing recommendations for the web application” paragraph in the Discussion; `streamlit_app/app.py` new expander block “Using data with different annotations or microarray data”.

R2.4 — Formatting (“Newomuscular” typo, gene italicisation)

There are a few formatting issues that need attention. Specifically, in the Supplementary Table 3 and 4 legends, some gene symbols appear to have formatting errors (e.g., “Newomuscular” instead of “Neuromuscular” in the main text). Additionally, please verify the consistency of the gene nomenclature (e.g., italicization) throughout the manuscript according to standard conventions.

Response: Thank you for catching these. We have performed a full pass over all gene-symbol occurrences in the main text, supplementary, and synchronised thesis files.

“Newomuscular” / Neuromuscular. On checking the current submission, all instances of “Neuromuscular” (Supplementary Table S3 module heading, Supplementary Table S4 module heading, Results Section 2.3, Discussion Section, Module 1 paragraph at line 205, and the concordance summary at line 211 of the original `paper.tex` file) are spelled correctly with no “Newomuscular” variant present. We suspect the reviewer may have read a PDF rendering with a font-rendering glitch in the production proof; we have nonetheless verified that the canonical source files (`paper/paper.tex`, `paper/supplementary.tex`) compile to PDFs in which every instance reads “Neuromuscular”. The re-compiled PDFs accompanying this response have been re-verified with `pdftotext` extraction plus `grep`.

Gene-name italicisation. HGNC/HGVS convention requires gene symbols to be in italic. We swept the manuscript, supplementary, and thesis sources and confirmed that all introduced gene symbols use `\textit{}`. One instance in the thesis Discussion that used plain-text gene names (“ACHE, NRK, KREMEN1, IGF2BP2, ...”) has been corrected to use `\textit{}`. We thank the reviewer for prompting this audit.

Manuscript changes: No source-text changes were necessary in `paper.tex` or `supplementary.tex` (already correctly formatted); `thesis/.../5_Discussion/5_Discussion.tex` gene-list sentence italicisation corrected.

R2.5 — Reference modernisation

The references are rather outdated and lack timeliness, especially the absence of literature from the past 3 to 5 years.

Response: We thank the reviewer for prompting this audit. The submitted manuscript already contained 36 references published in 2023–2026, but the audit identified clusters of older citations supporting newer claims, so we have strengthened the bibliography in several places.

Already-recent material strengthened during revision: the expanded cross-species methodology now cites the Cardoso-Moreira *et al.* 2019 atlas[1] and the ongoing dGTE project[2]; the batch-effect discussion cites the post-2023 PigGTE Compendium update [3, 4, 5]; the staging

biology discussion cites recent epigenetic and transcriptomic-clock work (Teschendorff 2025[6], Sun 2025[7], Muralidharan 2025[8], Lu 2023[9], Iwata 2024[10]); the marker-gene panel response (R3.3) cites 25 literature references underpinning the panel.

Newly proposed entries. We have written five additional 2022–2025 entries to `review/new_bib_entries.bib` for the corresponding author to paste into the shared bibliography file (per the project’s “bibliography-is-read-only” rule, the script does not modify `pig-age-human.bib` directly):

- Wu *et al.* 2025 (*BMC Biology*), single-cell pig skin atlas across 10 developmental stages—directly complementary porcine developmental resource.
- Bama pig single-cell atlas across seven tissues (*BMC Genomic Data* 2025).
- Comprehensive transcriptomic atlas of porcine immune tissues and PBMC dynamics (2025), relevant to our blood model.
- Tian *et al.* 2025 survey on ordinal regression, providing an up-to-date methodological reference for the Frank–Hall reduction.
- Wang *et al.* 2022 pig cell landscape (*Nature Communications*), the multi-tissue single-cell predecessor to which our bulk developmental atlas adds time-resolution.

Once the corresponding author has pasted these BibTeX entries into `pig-age-human.bib` (the file is curated manually to avoid hallucinated citations) we will cite them in the Background and Discussion—specifically (a) the pig skin atlas paper as a “contemporary porcine developmental resource at single-cell resolution” (Background, paragraph 3), (b) the ordinal-regression survey as a methodological reference (Methods, “Machine learning framework”), and (c) the Bama pig atlas and PBMC immune atlas as companion multi-tissue resources (Discussion, “species-agnostic framework”). The proposed insertion points are tracked in `review/new_bib_entries.bib` comments.

Manuscript changes: Five proposed 2022–2025 BibTeX entries appended to `review/new_bib_entries.bib`; `paper/paper.tex` citations will be added once the corresponding author has pasted the entries into `paper/pig-age-human/pig-age-human.bib` (we did not modify the bib file directly to comply with the project’s manual-curation rule that protects against hallucinated references).

Reviewer 3

R3.1 — Batch-effect PCA/UMAP across all tissues

Batch Effect Evaluation: The dataset integrates samples from the PigGTEx resource, which aggregates RNA-seq profiles from multiple sequencing platforms and library protocols. The authors argue that tree-based models are robust to monotonic platform shifts and provide within-project validation for muscle tissue. While this is a valid point, batch effects can be non-linear and potentially confounded with developmental stages, especially in tissues other than muscle. A simple visualization (e.g., UMAP or PCA) of the data colored by project or batch, rather than just by tissue, would strengthen the claim that batch effects do not drive the developmental predictions. Specifically, while the within-project validation for muscle is robust, similar validation or assessment for other tissues is missing. I recommend the authors provide an evaluation of batch effects across the entire dataset to demonstrate that the developmental signals are not artifacts of technical variation.

Response: We thank the reviewer for this important and well-targeted comment. We agree that within-project muscle validation alone is insufficient to demonstrate that batch effects are not driving developmental signals in the remaining tissues, and we have therefore added a global, multi-tissue batch-effect evaluation.

New analysis (Supplementary Fig. S4; Supplementary Table S6). For each of the five focal tissues (muscle, brain, liver, lung, blood) we $\log_2(\text{TPM}+1)$ -transformed the expression matrix, retained the 5,000 most-variable genes, and computed PCA (10 components) and UMAP (`n_neighbors=15`, `min_dist=0.1`, `metric="euclidean"`, `random_state=42`). For PC1–PC5 we fit OLS models $\text{PC} \sim \text{Stage}$, $\text{PC} \sim \text{BioProject}$, and the combined model with one-hot encoded factors, and report the corresponding R^2 values together with the silhouette scores in PC-10 space.

Honest framing of the results. Developmental stage alone explains a substantial fraction of PC1 variance in four of the five tissues—muscle $R^2_{\text{stage}} = 0.40$ (758 samples from 56 BioProjects), brain 0.65 (43/6), lung 0.33 (129/20), and blood 0.55 (78/7)—and the UMAPs in Supplementary Fig. S4 show visible infant-to-adult gradients within tissues even where multiple BioProjects co-occupy the same region of the embedding. Liver is the most batch-structured tissue at PC1 ($R^2_{\text{stage}} = 0.02$ vs $R^2_{\text{bioproject}} = 0.94$, 309 samples / 34 projects); its developmental component is recovered at PC5 ($R^2_{\text{stage}} = 0.29$), and held-out balanced accuracy for liver remains 0.921. We acknowledge in the Discussion that BioProject explains the larger absolute variance share in every tissue—an expected consequence of the PigGTEx sampling design, in which most projects sample a single narrow age window so that Stage and BioProject are partially confounded by design—and we therefore do not claim that the embedding is batch-free. The combined evidence (within-project muscle fold-change replication, visible UMAP stage gradients in 4/5 tissues, held-out balanced accuracies of 0.81–0.96, and the cross-species correlation of $r = 0.69$ in muscle) supports the interpretation that the developmental signal captured by our models is biological rather than purely technical. The full embeddings, variance-decomposition table, and silhouette scores are reproducible from `review/analyses/batch_effect_visualization.py`.

Manuscript changes: New Supplementary Fig. S4 and Supplementary Table S6; `paper/paper.tex` Discussion “Several limitations...” paragraph now references and quantifies the batch-effect analysis; `thesis/.../5_Discussion/5_Discussion.tex` “Robustness to Batch Effects” subsection updated to match.

R3.2 — Cross-species methodology (one-to-one orthologs)

Details on Cross-Species Comparison: The cross-species validation is a key strength of the paper, identifying 97% directional concordance in gene trajectories. However, the methodology for this comparison lacks sufficient detail. Specifically, how were orthologous genes between pig and human defined? Was a strict one-to-one orthologue mapping used? Additionally, while the authors used $|\log_2 \text{FC}|$ comparisons to avoid direct expression level comparisons, the pre-processing steps for the human dataset need to be more explicitly described to ensure reproducibility. Clarifying these steps is crucial for assessing the validity of the correlation observed ($r=0.69$).

Response: Yes, a strict one-to-one ortholog mapping is used; we have made every step explicit in the revised Methods and in the analysis script docstrings.

Ortholog filter. Pig (*Sus scrofa*) and human Ensembl gene identifiers were resolved to HGNC-style gene symbols via `mygene.info` (the resulting maps are cached at `data/cardoso_moreira_2019/_pig_id_to_` and `_human_id_to_symbol.json`). We retained only symbols appearing *exactly once* in each species’ annotation table (“shared-unique-symbol” filter), and then *verified* each candidate pair against the Ensembl Compara REST homology endpoint (`/homology/id/sus_scrofa/{ENSSSCG_id}` with `target_species=human`, `type=orthologues`), keeping only pairs whose reciprocal homology classification was `ortholog_one2one` (the Ensembl Compara confidence-1 equivalent for strict one-to-one orthologues). The full filtered table is cached at `review/analyses/results/pig_human_one_to_one_orth` and is used by all new cross-species analyses.

Human preprocessing pipeline. The Cardoso-Moreira 2019 human preprocessing is now

documented step-by-step in Methods and in the script docstring of `review/analyses/cross_species_all_tissue`

- (i) load `Human.RPKM.tsv` (derived from `Human.RPKM.txt` as distributed under E-MTAB-6814);
- (ii) restrict to columns matching the pattern `<Tissue>.<stage>.<n>` for each tissue of interest;
- (iii) drop any column whose stage label ends in `wpc` to retain postnatal samples only; (iv) map Cardoso-Moreira postnatal stage names to the porcine five-stage scheme (newborn/infant/toddler $\rightarrow$ Infant; school/youngTeenager $\rightarrow$ Pre-pubertal; teenager/oldTeenager $\rightarrow$ Post-pubertal; youngAdult/youngMidAge $\rightarrow$ Adult; olderMidAge/senior $\rightarrow$ Adult); (v) compute $\log_2$ fold-change between Infant and Adult with a pseudocount of 0.01 on raw RPKM (matching the pig TPM pipeline) and apply Welch’s *t*-test per gene (more powerful at the small per-tissue postnatal *n* in this atlas); (vi) Benjamini–Hochberg FDR correction; (vii) restrict to the validated one-to-one ortholog set. The pig pipeline is symmetric (TPM, Mann–Whitney *U* for the larger sample sizes, identical FDR and ortholog filter). The original muscle analysis (Schaiter *et al.*, *n* = 11) used precisely the same pipeline, simply restricted to the muscle expression matrix.

Manuscript changes: **paper/paper.tex** Methods “Cross-species evolutionary analysis” subsection now describes the ortholog filter and human preprocessing pipeline; **review/analyses/cross_species_all_tissues.py** docstring is the canonical reproducible reference; **review/analyses/results/pig_human_one_to_one_orthologs.csv** contains the validated mapping.

R3.3 — Biological validation of stage boundaries

Biological Validation of Developmental Stages: The developmental stages are defined based on physiological milestones (infant, early childhood, etc.). The machine learning model predicts these stages with high accuracy. However, this only proves that the stages are transcriptionally distinct, not that the boundaries defined by chronological age correspond precisely to biological maturation switches. To further validate these stage definitions, the authors could analyze the expression of known developmental marker genes or pathways at the transition points between the defined stages. For instance, showing that genes related to weaning or puberty show sharp expression changes around the defined boundaries would provide independent biological support for the staging framework.

Response: We thank the reviewer for this excellent suggestion. To provide independent biological support for the staging framework, we have added a curated marker-gene validation (Supplementary Fig. S6; `review/analyses/results/marker_gene_transitions.csv`).

Panel. 25 literature-supported postnatal mammalian markers were selected to span the four transition windows of our framework: suckling $\rightarrow$ weaning at $\sim$ 21 d (*AFP*, *ALB*, *IGF1*, *IGF2*, *LCT*, ...); weaning $\rightarrow$ growth at $\sim$ 60 d (*MYH3*, *MYH8*, *MYH2*, *MSTN*, *IGFBP5*, *MBP*, *PLP1*, *SYP*); the pubertal switch at $\sim$ 150 d (*INSL3*, *LHB*, *AR*, *ESR1*, *GHR*); and adult maturation at $\sim$ 365 d (*CDKN2A*, *LMNA*, *SIRT1*, *SIRT3*, *MKI67*).

Test. For each marker $\times$ tissue pair we computed mean $\log_2(\text{TPM} + 1)$ per stage, fit one-way ANOVA with Tukey–HSD post-hoc, and recorded the adjacent-stage pair with the largest mean difference. After excluding pairs with fewer than five samples per flanking stage, 30 marker $\times$ tissue combinations were evaluable.

Results. 20/30 (67%) matched the literature-predicted direction of change; 10/30 (33%) had their largest single-step change at the exact predicted transition window; T1 (suckling $\rightarrow$ weaning) had the highest window-match rate ($4/7 = 57\%$). The most informative miss is the porcine myosin heavy-chain switch (*MYH3*, *MYH8*, *MYH2*), which was directionally correct but transitioned earlier than the assumed boundary (Infant $\rightarrow$ Early childhood rather than Early childhood $\rightarrow$ Pre-pubertal), consistent with the first-month perinatal isoform exchange reported in pig skeletal muscle.

We have not used these results to retro-fit the stage boundaries (doing so would invalidate the model evaluation that we report); rather we have added a paragraph in the Discussion presenting this analysis as *partial*, *independent* biological support for the framework together

with an explicit qualification that some boundaries (notably the myosin-switch midpoint and the post-pubertal→adult boundary in lung and brain, which are sparsely sampled) are approximations. Twenty-five proposed BibTeX entries for the panel references have been written to `review/new_bib_entries.bib` so the corresponding author can paste them into the shared bibliography.

Manuscript changes: New Supplementary Fig. S6; `paper/paper.tex` Discussion paragraph immediately following the “noisy-labels” paragraph; `thesis/.../5_Discussion/5_Discussion.tex` new “Independent Marker-Gene Validation of Stage Boundaries” subsection.

R3.4 — Tissue-specific top features

Interpretability of Tissue-Specific Models: The paper highlights that developmental programs are profoundly tissue-specific, with extremely low cross-tissue feature overlap. While the cross-species markers are discussed in detail, the tissue-specific features driving the predictions in the brain, liver, lung, and blood are less explored. Providing the top feature genes for each tissue-specific model and their biological relevance would significantly enhance the utility of the atlas, offering insights into the unique molecular logic of each tissue’s development.

Response: We thank the reviewer for this constructive suggestion; addressing it has materially improved the interpretability of the atlas. We have added Supplementary Table S7 listing the top 20 features per tissue ranked by aggregate LightGBM importance, together with a brief functional description (drawn from mygene.info and UniProt REST, with placeholder rows for top-ranked features that returned no curated description from either resource—predominantly novel pig loci and uncharacterised LOC entries), the peak developmental stage, and the peak-vs-trough fold change in mean $\log_2(\text{TPM} + 1)$. A companion heatmap (Supplementary Fig. S5) displays row-z-scored trajectories of the top 10 genes per tissue across the five stages, making the timing of each transition immediately visible.

For the four non-muscle tissues, biologically coherent stage-defining genes that were not previously discussed emerge:

- **Brain:** ketogenic enzyme *HMGCS2* (infant peak) and Hippo-pathway kinase *LATS1* (post-pubertal peak), consistent with a metabolic-to-circuit-maturation transition during postnatal cortical development.
- **Liver:** *ALDH18A1* (proline/ornithine biosynthesis) and *COL1A1* peak in infancy alongside post-pubertal-peaking transcription factors *SATB2* and *CDH7*, mirroring the amino-acid-anabolism-to-mature-hepatocyte programme.
- **Lung:** matricellular *CCN5* (post-pubertal; fold change 2.71) and adhesion molecule *L1CAM* dominate alveolar maturation; the oncofetal RNA-binder *IGF2BP2* (infant peak)—also a top cross-species muscle marker—signals neonatal mesenchymal growth, suggesting a shared early-postnatal proliferation programme across mesodermal tissues.
- **Blood:** ribosomal *RPLP1* (infant peak; fold change 3.14) and the chemokine receptor *CCR6* (post-pubertal) track a neonatal ribosomal/transcriptional surge and adaptive-immune maturation respectively.

Placeholder rows in Supplementary Table S7 are retained intentionally so that readers can see which top-ranked features remain functionally uncharacterised in the pig genome—identifying clear opportunities for follow-up annotation. The table and figure are reproducible from `review/analyses/tissue_specific_features.py` (the script caches API responses locally).

Manuscript changes: New Supplementary Fig. S5 and Supplementary Table S7; `paper/paper.tex` Results section 2.3 has a new paragraph after the cross-tissue feature-overlap discussion; `thesis/.../4_Results/4_Results.tex` new “Tissue-Specific Top Features” subsection added to Section 4.

Minor concerns

R3.m1 — Confusion matrices reference

Performance Metrics: The performance metrics provided are comprehensive. However, the confusion matrices and precision/recall heatmaps could be better integrated into the main text or referred to more explicitly to highlight specific classification challenges (e.g., misclassification between adjacent stages).

Response: Agreed. We have rewritten the paragraph that introduces the confusion matrices in Results 2.2 to surface the matrices more explicitly and to tie them to the per-class precision/recall/ F_1 heatmap (Supplementary Fig. S3a) and per-class values in Supplementary Table S2. The new text reads (paraphrased): “Inspection of the per-tissue confusion matrices (Fig. 2b) confirms that residual misclassifications are concentrated almost exclusively between *adjacent* developmental stages—for example, infant samples occasionally classified as early childhood, or pre-pubertal samples classified as post-pubertal—rather than being distributed randomly across the ordinal scale. Together with the per-stage precision/recall/ F_1 heatmap (Supplementary Fig. S3a) and per-class values reported in Supplementary Table S2, these patterns indicate that the most difficult decision boundaries lie at the genuinely continuous transitions of postnatal biology (notably the early-childhood-to-pre-pubertal transition, where post-weaning physiology overlaps with juvenile growth), rather than at boundaries of equal biological difficulty.” The figure and table references now appear in the same sentence as the narrative claim, addressing the reviewer’s concern.

Manuscript changes: ***paper/paper.tex*** Results 2.2 paragraph immediately following the ordinal-consistency sentence.

R3.m2 — Infant-vs-adult comparison gap implications

Limitations: The authors acknowledge the limitations of the human dataset (small sample size, lack of intermediate stages). It would be beneficial to further discuss the implications of the “infant vs. adult” comparison gap. How might the inclusion of childhood/adolescent data affect the observed correlations? Addressing this theoretically would strengthen the discussion.

Response: We have expanded the relevant Limitations paragraph to address this directly. The text now spells out that the reported $r = 0.69$ summarises agreement of *net* infant-to-adult $\log_2$ fold changes, not of trajectory shapes, and enumerates three possible outcomes had childhood/adolescent samples been available for any given gene:

- the gene’s monotonic infant→adult direction is preserved through intermediate ages, strengthening the conservation claim for that gene;
- the gene exhibits a transient peak or trough at an intermediate human stage (e.g. around puberty), refining—rather than overturning—the directional concordance and improving pig↔human stage alignment;
- the trajectory shape diverges qualitatively, identifying species-specific re-routing of the developmental program.

Because our current analysis cannot distinguish these scenarios, we have framed all conservation claims in *net* infant-to-adult terms and note that our cross-species inferences should be read as evidence of *directional* conservation of the overall maturation programme rather than of *trajectory-shape* conservation. Resolving this will require human postnatal muscle RNA-seq cohorts with at least one childhood and one adolescent sample group; we explicitly flag this as a priority for future work and point at dGTEx [2] and the multi-organ developmental atlas of

Cardoso-Moreira *et al.* [1] as the natural data sources.

Manuscript changes: `paper/paper.tex` expanded “Third...” point of the Limitations paragraph in the Discussion; `thesis/.../5_Discussion/5_Discussion.tex` expanded “Human Validation Dataset Size” subsubsection.

Closing

We are grateful to the three reviewers for the time and care they have invested in our manuscript. The revisions have, we believe, substantially strengthened both the methodological rigour and the translational reach of the study. We are very happy to provide any further clarifications the reviewers or editors may require.

Sincerely,

Peter S. Theobald (corresponding author), on behalf of all co-authors